

# Chem277B: Machine Learning Algorithms

## Homework assignment #7: Deeper Learning and Regularization

```
In [44]: import numpy as np
import pandas as pd
import math
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
from torch.optim import SGD, Adam
from sklearn.model_selection import train_test_split, KFold
from sklearn import cluster, datasets, mixture
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from itertools import cycle, islice
from pylab import *
import seaborn as sns
from functools import wraps
from time import time
import random
from tqdm import tqdm
import warnings

sns.set()
```

### 1. Bias-variance tradeoff.

(a) I have sorted out the training and testing datasets and normalized the data using each  $32 \times 32$  image's maximum pixel value. I noticed some of the image's max pixel value is not 255. I also did a little confirmation test by printing out the sum of the training/testing sets before and after normalization. There's indeed a >200-fold decrease of the sum value.

```
In [2]: # First load the mnist data and convert all elements into lists / arrays
mnist = list(pd.read_pickle('mnist.pkl'))
mnist[0] = list(mnist[0])
mnist[1] = list(mnist[1])
train_X = mnist[0][0]
train_y = mnist[0][1]
```

```
test_X = mnist[1][0]
test_y = mnist[1][1]
print('Train: X=%s, y=%s' % (train_X.shape, train_y.shape))
print('Test: X=%s, y=%s' % (test_X.shape, test_y.shape))
```

```
Train: X=(60000, 32, 32), y=(60000,)
Test: X=(10000, 32, 32), y=(10000,)
```

```
In [3]: # Define a function to normalize the training and testing data sets
def normalize_pixels(train_X, test_X):

    # First convert the dataset to floats
    train_X_norm = train_X.astype('float32')
    test_X_norm = test_X.astype('float32')

    # Find maximum values for all 60000 / 10000 pictures in train and test datasets
    # and broadcast to a (60000 / 10000, 32, 32) shape array
    train_X_max = np.broadcast_to(train_X.max(axis=(1,2))[:, np.newaxis, np.newaxis], (60000, 32, 32))
    test_X_max = np.broadcast_to(test_X.max(axis=(1,2))[:, np.newaxis, np.newaxis], (10000, 32, 32))

    # Normalize the datasets
    train_X_norm = train_X_norm / train_X_max
    test_X_norm = test_X_norm / test_X_max

    # Return the normalized datasets
    return train_X_norm, test_X_norm
```

```
In [4]: # Normalize the datasets and confirm the data has been normalized
train_X_norm, test_X_norm = normalize_pixels(train_X, test_X)
print(train_X.sum(), train_X_norm.sum())
print(test_X.sum(), test_X_norm.sum())
```

```
1567298545 6148662.5
264923200 1039329.2
```

(b) The finished Trainer class and the ANN are shown below.

I divided the training data into a 3-fold groups of training and validation datasets using KFold modality. I chose to do the training-validation split outside the Trainer Class because I found that doing it inside the class always creates index errors and it's very hard to debug. The index errors happened because the for loop to run epoch training first shuffled the indices so some of the large indices will appear and they are outside the index range. I haven't found an elegant way to solve the problem. So I just put the KFold split outside the Trainer Class and it worked just fine.

```
In [40]: def timing(f):
    @wraps(f)
    def wrap(*args, **kw):
        ts = time()
        result = f(*args, **kw)
```

```

    te = time()
    print('func:%r  took: %2.4f sec' % (f.__name__, te-ts))
    return result
return wrap

def create_chunks(complete_list, chunk_size=None, num_chunks=None):
    """
    Cut a list into multiple chunks, each having chunk_size (the last chunk might be less than chunk_size)
    or having a total of num_chunk chunks
    """
    chunks = []
    if num_chunks is None:
        num_chunks = math.ceil(len(complete_list) / chunk_size)
    elif chunk_size is None:
        chunk_size = math.ceil(len(complete_list) / num_chunks)
    for i in range(num_chunks):
        chunks.append(complete_list[i * chunk_size: (i + 1) * chunk_size])
    return chunks

class Trainer():
    def __init__(self, model, optimizer_type, learning_rate, epoch, batch_size, input_transform=lambda x: x):
        """ The class for training the model
        model: nn.Module
            A pytorch model
        optimizer_type: 'adam' or 'sgd'
        learning_rate: float
        epoch: int
        batch_size: int
        input_transform: func
            transforming input. Can do reshape here
        """
        self.model = model
        if optimizer_type == "sgd":
            self.optimizer = SGD(model.parameters(), learning_rate, momentum=0.9)
        elif optimizer_type == "adam":
            self.optimizer = optim.Adam(model.parameters(), learning_rate)

        self.epoch = epoch
        self.batch_size = batch_size
        self.input_transform = input_transform

    @timing
    def train(self, inputs, outputs, val_inputs, val_outputs, early_stop=False, l2=False, silent=False):
        """ train self.model with specified arguments using 3-fold cross-validation
        inputs: np.array, The shape of input_transform(input) should be (ndata,nfeatures)
        outputs: np.array shape (ndata,)
        val_nputs: np.array, The shape of input_transform(val_input) should be (ndata,nfeatures)
        val_outputs: np.array shape (ndata,)
        early_stop: bool
        l2: bool

```

`silent: bool.` Controls whether or not to print the train and val error during training

`@return`

a dictionary of arrays with train and val losses and accuracies

"""

### convert data to tensor of correct shape and type here ###

`inputs = torch.Tensor(self.input_transform(inputs)).float().clone().detach()` # inputs are (X (40000), 32, 32))

`outputs = torch.Tensor(outputs).long().clone().detach()` # outputs are (y (40000), )

`val_inputs = torch.Tensor(self.input_transform(val_inputs)).float().clone().detach()` # val\_inputs are (X (20000), 32, 32))

`val_outputs = torch.Tensor(val_outputs).long().clone().detach()` # val\_outputs are (y (20000), )

`losses = []`

`accuracies = []`

`val_losses = []`

`val_accuracies = []`

`weights = self.model.state_dict()`

`lowest_val_loss = np.inf`

**for** `n_epoch` **in** `tqdm(range(self.epoch), leave=False):`

`self.model.train()`

`batch_indices = list(range(inputs.shape[0]))` # range(40000)

`random.shuffle(batch_indices)`

`batch_indices = create_chunks(batch_indices, chunk_size=self.batch_size)`

`epoch_loss = 0`

`epoch_acc = 0`

**for** `batch` **in** `batch_indices:`

`batch_importance = len(batch) / len(outputs)`

`batch_input = inputs[batch]`

`batch_output = outputs[batch]`

    ### make prediction and compute loss with loss function of your choice on this batch ###

`batch_predictions = self.model(batch_input)`

`loss = F.cross_entropy(batch_predictions, batch_output)`

**if** `l2:`

        ### Compute the loss with L2 regularization ###

`l2_lambda = 1e-5`

`l2_reg = torch.tensor(0.)`

**for** `param` **in** `self.model.parameters():`

`l2_reg += torch.norm(param)`

`loss += l2_lambda * l2_reg`

`self.optimizer.zero_grad()`

`loss.backward()`

`self.optimizer.step()`

    ### Compute epoch\_loss and epoch\_acc

`epoch_loss += loss.item() * batch_importance`

`epoch_acc += (batch_predictions.argmax(dim=1) == batch_output).float().mean().item() * batch_importance`

```

        val_loss, val_acc = self.evaluate(val_inputs, val_outputs, print_acc=False)

        if n_epoch % 10 == 0 and not silent:
            print("Epoch %d/%d - Loss: %.3f - Acc: %.3f" % (n_epoch + 1, self.epoch, epoch_loss, epoch_acc))
            print("                Val_loss: %.3f - Val_acc: %.3f" % (val_loss, val_acc))
            losses.append(epoch_loss)
            accuracies.append(epoch_acc)
            val_losses.append(val_loss)
            val_accuracies.append(val_acc)
            if early_stop:
                if val_loss < lowest_val_loss:
                    lowest_val_loss = val_loss
                    weights = self.model.state_dict()

        if early_stop:
            self.model.load_state_dict(weights)

    return {"losses": losses, "accuracies": accuracies, "val_losses": val_losses, "val_accuracies": val_accuracies}

def evaluate(self, inputs, outputs, print_acc=True):
    """ evaluate model on provided input and output
    inputs: np.array, The shape of input_transform(input) should be (ndata,nfeatures)
    outputs: np.array shape (ndata,)
    print_acc: bool

    @return
    losses: float
    acc: float
    """
    with torch.no_grad():
        inputs = torch.tensor(inputs).float().clone().detach()
        outputs = torch.tensor(outputs).long().clone().detach()
        outputs_pred = self.model(inputs)
        loss = F.cross_entropy(outputs_pred, outputs)
        acc = torch.mean((torch.argmax(outputs_pred, dim=1) == outputs).float())
    if print_acc:
        print(f'Accuracy: {acc.item()}')
    return loss, acc

```

In [29]:

```

class MLP(nn.Module):
    def __init__(self):
        super(MLP, self).__init__()
        self.layers = nn.Sequential(
            nn.Flatten(),
            nn.Linear(1024, 3),
            nn.Sigmoid(),
            nn.Linear(3, 10),
            nn.Sigmoid()

```

```
)
def forward(self, X):
    return self.layers(X)
```

(c) I made the multi-layer perceptron ANN with 2 computing layers, a hidden layer of 3 neurons and a final output layer of 10 neurons. Both layers use a sigmoid activation function.

I then ran the model using the required ADAM optimizer and parameters. From the results and the generated plots, clearly the model converges without the need or further regularization. But the accuracies are pretty bad and stay at ~0.58. I tested the model on the testing datasets and got similar results on loss and accuracy.

The bias-variance tradeoff refers to the balancing of the model's ability to fit the training data (low bias) with its ability to generalize to new data (low variance).

The MLP ANN we employ here consists of two linear layers with a non-linear activation function in between. The first linear layer maps the input data from a 1024-dimensional space to a 3-dimensional space. The second linear layer maps the output of the first layer from a 3-dimensional space to a 10-dimensional space. The mapping of the space dimension itself is not reasonable as indicated in the tutorial class, gradient vanishing could happen. The hidden layer with 3-dimensions may not provide enough capacity for the model to learn complex patterns in the data.

Hence I think the model's the poor accuracy is due to high bias and limited by the hidden layer and the activation function. The model's variance seems OK but it's limited by the bias.

```
In [30]: mlp = MLP()
ann_1_adam = Trainer(mlp, "adam", 2e-3, 50, 128, input_transform=lambda x: x,)

# Split the dataset into 3-fold training and validation sets
training_result_all = []
kf = KFold(n_splits=3, shuffle=True, random_state=1)
for train_index, val_index in kf.split(train_X_norm, train_y):
    X_train, X_val = train_X_norm[train_index], train_X_norm[val_index]
    y_train, y_val = train_y[train_index], train_y[val_index]

    training_result = ann_1_adam.train(X_train, y_train, X_val, y_val, \
                                       early_stop=False, l2=False, silent=False)
    training_result_all.append(training_result)
```

```
0%|          | 0/50 [00:00<?, ?it/s]<ipython-input-8-c6373bc9576a>:137: UserWarning: To copy construct from a tensor,
it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than
torch.tensor(sourceTensor).
  inputs = torch.tensor(inputs).float().clone().detach()
<ipython-input-8-c6373bc9576a>:138: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor
r.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).
  outputs = torch.tensor(outputs).long().clone().detach()
2%||         | 1/50 [00:00<00:38, 1.29it/s]
Epoch 1/50 - Loss: 2.206 - Acc: 0.268
              Val_loss: 2.133 - Val_acc: 0.318

22%|█        | 11/50 [00:06<00:20, 1.90it/s]
Epoch 11/50 - Loss: 1.790 - Acc: 0.535
              Val_loss: 1.792 - Val_acc: 0.527

42%|██       | 21/50 [00:11<00:14, 2.04it/s]
Epoch 21/50 - Loss: 1.739 - Acc: 0.504
              Val_loss: 1.748 - Val_acc: 0.490

62%|████    | 31/50 [00:16<00:11, 1.70it/s]
Epoch 31/50 - Loss: 1.722 - Acc: 0.495
              Val_loss: 1.734 - Val_acc: 0.480

82%|██████  | 41/50 [00:23<00:05, 1.68it/s]
Epoch 41/50 - Loss: 1.707 - Acc: 0.483
              Val_loss: 1.719 - Val_acc: 0.465

func:'train' took: 29.6902 sec

2%||         | 1/50 [00:00<00:23, 2.07it/s]
Epoch 1/50 - Loss: 1.705 - Acc: 0.473
              Val_loss: 1.699 - Val_acc: 0.485

22%|█        | 11/50 [00:06<00:23, 1.68it/s]
Epoch 11/50 - Loss: 1.696 - Acc: 0.476
              Val_loss: 1.700 - Val_acc: 0.479

42%|██       | 21/50 [00:13<00:25, 1.14it/s]
Epoch 21/50 - Loss: 1.691 - Acc: 0.477
              Val_loss: 1.699 - Val_acc: 0.479

62%|████    | 31/50 [00:19<00:11, 1.60it/s]
Epoch 31/50 - Loss: 1.686 - Acc: 0.479
              Val_loss: 1.696 - Val_acc: 0.480

82%|██████  | 41/50 [00:25<00:05, 1.71it/s]
Epoch 41/50 - Loss: 1.670 - Acc: 0.480
              Val_loss: 1.682 - Val_acc: 0.480

func:'train' took: 32.0851 sec

2%||         | 1/50 [00:00<00:23, 2.09it/s]
Epoch 1/50 - Loss: 1.672 - Acc: 0.483
              Val_loss: 1.661 - Val_acc: 0.489

22%|█        | 11/50 [00:06<00:25, 1.56it/s]
Epoch 11/50 - Loss: 1.664 - Acc: 0.506
              Val_loss: 1.664 - Val_acc: 0.505
```

```
42%|███████| 21/50 [00:11<00:14, 1.97it/s]
Epoch 21/50 - Loss: 1.660 - Acc: 0.569
              Val_loss: 1.663 - Val_acc: 0.569

62%|███████| 31/50 [00:18<00:12, 1.52it/s]
Epoch 31/50 - Loss: 1.657 - Acc: 0.580
              Val_loss: 1.663 - Val_acc: 0.580

82%|███████| 41/50 [00:23<00:04, 2.02it/s]
Epoch 41/50 - Loss: 1.655 - Acc: 0.583
              Val_loss: 1.664 - Val_acc: 0.579

func:'train' took: 28.4268 sec
```


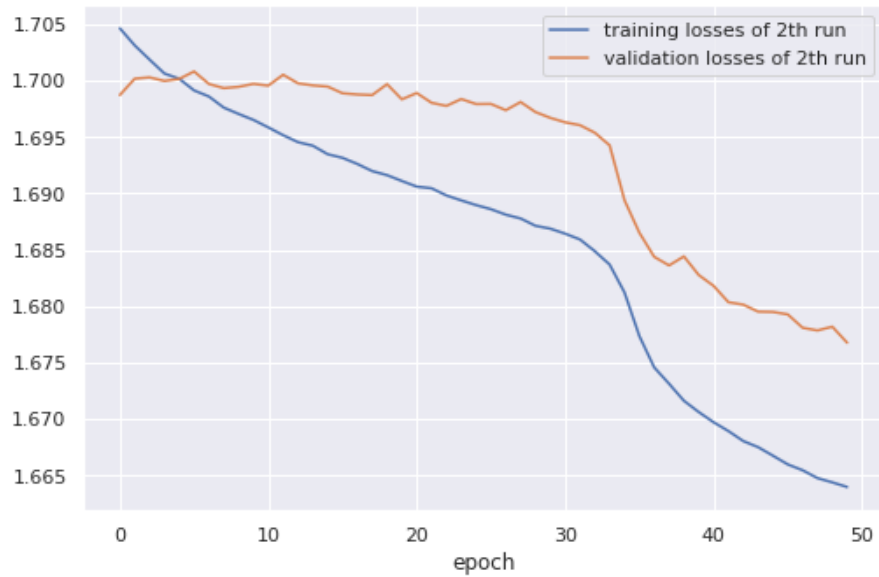


```
In [31]: ann_1_adam.evaluate(test_X_norm, test_y)
```

```
Accuracy: 0.5806999802589417
```

```
Out[31]: (tensor(1.6666), tensor(0.5807))
```

```
In [32]: fig, axes = plt.subplots(3, 2, figsize=(18,18))
         for i in range(len(training_result_all)):
             axes[i][0].plot(training_result_all[i]["losses"], label = f"training losses of {i+1}th run")
             axes[i][0].plot(training_result_all[i]["val_losses"], label = f"validation losses of {i+1}th run")
             axes[i][0].legend()
             axes[i][0].set_xlabel('epoch')
             axes[i][1].plot(training_result_all[i]["accuracies"], label = f"training accuracies of {i+1}th run")
             axes[i][1].plot(training_result_all[i]["val_accuracies"], label = f"validation accuracies of {i+1}th run")
             axes[i][1].legend()
             axes[i][1].set_xlabel('epoch')
```





(d) After increasing the hidden layer size from 3 to 50, I noticed 2 big differences: 1) the training accuracies increased a lot from 0.5 to 0.98; 2) the training time is much longer, increasing from ~30 seconds to ~40 seconds.

In total the bias of the new model is significantly lower and it applies to the testing dataset very well too.

From what I read, the hidden layer size can affect the bias because it determines the number of parameters that the model can learn. In general a larger hidden layer size allows the model to learn more complex representations of the input data. Hence it allows the ANN model to capture patterns and relationships in the data. But the bias-variance tradeoff could start to happen when the hidden layer size is too large and leads to overfitting.

In this case, after increasing the hidden layer size from 3 to 50, both bias and variance improve on training, validation and testing datasets. I haven't found any indication of overfitting yet.

```
In [33]: class MLP2(nn.Module):
        def __init__(self):
            super(MLP2, self).__init__()
            self.layers = nn.Sequential(
                nn.Flatten(),
                nn.Linear(1024, 50),
                nn.Sigmoid(),
                nn.Linear(50, 10),
                nn.Sigmoid()
            )

        def forward(self, X):
            return self.layers(X)
```

```
In [34]: mlp2 = MLP2()
ann_2_adam = Trainer(mlp2, "adam", 2e-3, 50, 128, input_transform=lambda x: x,)

# Split the dataset into 3-fold training and validation sets
training_result_all = []
kf = KFold(n_splits=3, shuffle=True, random_state=1)
for train_index, val_index in kf.split(train_X_norm, train_y):
    X_train, X_val = train_X_norm[train_index], train_X_norm[val_index]
    y_train, y_val = train_y[train_index], train_y[val_index]

    training_result = ann_2_adam.train(X_train, y_train, X_val, y_val, \
                                       early_stop=False, l2=False, silent=False)
    training_result_all.append(training_result)
```

```
0%|          | 0/50 [00:00<?, ?it/s]<ipython-input-8-c6373bc9576a>:137: UserWarning: To copy construct from a tensor,
it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather th
an torch.tensor(sourceTensor).
  inputs = torch.tensor(inputs).float().clone().detach()
<ipython-input-8-c6373bc9576a>:138: UserWarning: To copy construct from a tensor, it is recommended to use sourceTens
r.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).
  outputs = torch.tensor(outputs).long().clone().detach()
2%||         | 1/50 [00:01<01:06, 1.36s/it]
Epoch 1/50 - Loss: 1.824 - Acc: 0.793
              Val_loss: 1.649 - Val_acc: 0.888
22%|██       | 11/50 [00:09<00:33, 1.17it/s]
Epoch 11/50 - Loss: 1.511 - Acc: 0.952
              Val_loss: 1.521 - Val_acc: 0.939
42%|████      | 21/50 [00:17<00:22, 1.28it/s]
Epoch 21/50 - Loss: 1.493 - Acc: 0.968
              Val_loss: 1.511 - Val_acc: 0.948
62%|██████    | 31/50 [00:26<00:14, 1.29it/s]
Epoch 31/50 - Loss: 1.484 - Acc: 0.975
              Val_loss: 1.508 - Val_acc: 0.953
82%|████████   | 41/50 [00:34<00:07, 1.18it/s]
Epoch 41/50 - Loss: 1.479 - Acc: 0.980
              Val_loss: 1.508 - Val_acc: 0.954
func:'train' took: 41.5278 sec
2%||         | 1/50 [00:00<00:37, 1.32it/s]
Epoch 1/50 - Loss: 1.492 - Acc: 0.968
              Val_loss: 1.481 - Val_acc: 0.979
22%|██       | 11/50 [00:09<00:30, 1.29it/s]
Epoch 11/50 - Loss: 1.479 - Acc: 0.979
              Val_loss: 1.484 - Val_acc: 0.977
42%|████      | 21/50 [00:18<00:25, 1.13it/s]
Epoch 21/50 - Loss: 1.476 - Acc: 0.982
              Val_loss: 1.486 - Val_acc: 0.975
62%|██████    | 31/50 [00:28<00:20, 1.09s/it]
Epoch 31/50 - Loss: 1.474 - Acc: 0.984
              Val_loss: 1.489 - Val_acc: 0.973
82%|████████   | 41/50 [00:36<00:07, 1.28it/s]
Epoch 41/50 - Loss: 1.473 - Acc: 0.985
              Val_loss: 1.491 - Val_acc: 0.972
func:'train' took: 45.3296 sec
2%||         | 1/50 [00:00<00:37, 1.32it/s]
Epoch 1/50 - Loss: 1.484 - Acc: 0.977
              Val_loss: 1.475 - Val_acc: 0.986
22%|██       | 11/50 [00:09<00:35, 1.10it/s]
```

```
Epoch 11/50 - Loss: 1.475 - Acc: 0.983
              Val_loss: 1.474 - Val_acc: 0.987
42%|██████    | 21/50 [00:17<00:22, 1.28it/s]
Epoch 21/50 - Loss: 1.473 - Acc: 0.984
              Val_loss: 1.476 - Val_acc: 0.986
62%|███████   | 31/50 [00:25<00:15, 1.22it/s]
Epoch 31/50 - Loss: 1.473 - Acc: 0.985
              Val_loss: 1.478 - Val_acc: 0.984
82%|█████████ | 41/50 [00:34<00:08, 1.08it/s]
Epoch 41/50 - Loss: 1.472 - Acc: 0.986
              Val_loss: 1.480 - Val_acc: 0.983
```


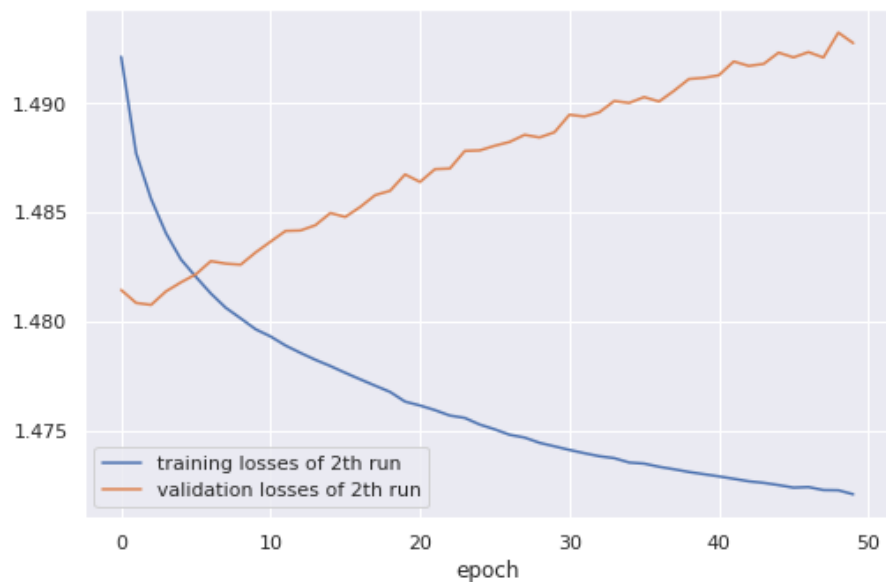


```
func:'train' took: 42.0510 sec
```

```
In [35]: ann_2_adam.evaluate(test_X_norm, test_y)
```

```
Accuracy: 0.9611999988555908
```

```
Out[35]: (tensor(1.5040), tensor(0.9612))
```

```
In [36]: fig, axes = plt.subplots(3, 2, figsize=(18,18))
for i in range(len(training_result_all)):
    axes[i][0].plot(training_result_all[i]["losses"], label = f"training losses of {i+1}th run")
    axes[i][0].plot(training_result_all[i]["val_losses"], label = f"validation losses of {i+1}th run")
    axes[i][0].legend()
    axes[i][0].set_xlabel('epoch')
    axes[i][1].plot(training_result_all[i]["accuracies"], label = f"training accuracies of {i+1}th run")
    axes[i][1].plot(training_result_all[i]["val_accuracies"], label = f"validation accuracies of {i+1}th run")
    axes[i][1].legend()
    axes[i][1].set_xlabel('epoch')
```



## 2. Deep Learning and regularization.

(a) I modified the ANN and added a dropout layer with 15% probability after the hidden layer. From what I read and the tutorial this is to prevent overfitting. The end result is actually worse than the ANN without the dropout layer. I have tabulated the training / validation / test accuracies below.

Category	Training Accuracy	Validation Accuracy	Test Accuracy	Run Time(seconds)
1c: Hidden Layer=3	0.583	0.579	0.581	28.4
1d: Hidden Layer=50	0.986	0.983	0.961	42.1
2a: 1d + 15% dropout	0.976	0.970	0.956	43.6

From the comparison, it's very clear that 15% dropout didn't really help the accuracy and the application to test accuracy. Which is consistent with 1d's observation that the ANN is not overfitting yet. The run time actually increased a little likely due to the random dropping step.

```
In [37]: class MLP2_dropout(nn.Module):
        def __init__(self):
            super(MLP2_dropout, self).__init__()
            self.layers = nn.Sequential(
                nn.Flatten(),
                nn.Linear(1024, 50),
                nn.Sigmoid(),

                # Add a dropout layer with 15% probability
                nn.Dropout(p=0.15),

                nn.Linear(50, 10),
                nn.Sigmoid()
            )

        def forward(self, X):
            return self.layers(X)
```

```
In [38]: mlp2_dropout = MLP2_dropout()
ann_3_adam = Trainer(mlp2_dropout, "adam", 2e-3, 50, 128, input_transform=lambda x: x,)
```

```
# Split the dataset into 3-fold training and validation sets
```

```
training_result_all = []  
kf = KFold(n_splits=3, shuffle=True, random_state=1)  
for train_index, val_index in kf.split(train_X_norm, train_y):  
    X_train, X_val = train_X_norm[train_index], train_X_norm[val_index]  
    y_train, y_val = train_y[train_index], train_y[val_index]  
  
    training_result = ann_3_adam.train(X_train, y_train, X_val, y_val, \  
        early_stop=False, l2=False, silent=False)  
    training_result_all.append(training_result)
```

```
0%|          | 0/50 [00:00<?, ?it/s]<ipython-input-8-c6373bc9576a>:137: UserWarning: To copy construct from a tensor,  
it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather th  
an torch.tensor(sourceTensor).
```

```
    inputs = torch.tensor(inputs).float().clone().detach()
```

```
<ipython-input-8-c6373bc9576a>:138: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor  
.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).
```

```
    outputs = torch.tensor(outputs).long().clone().detach()
```

```
2%||          | 1/50 [00:01<00:50, 1.03s/it]
```

```
Epoch 1/50 - Loss: 1.839 - Acc: 0.777
```

```
Val_loss: 1.662 - Val_acc: 0.874
```

```
22%|██        | 11/50 [00:09<00:31, 1.25it/s]
```

```
Epoch 11/50 - Loss: 1.522 - Acc: 0.943
```

```
Val_loss: 1.530 - Val_acc: 0.934
```

```
42%|██████     | 21/50 [00:18<00:27, 1.05it/s]
```

```
Epoch 21/50 - Loss: 1.506 - Acc: 0.957
```

```
Val_loss: 1.519 - Val_acc: 0.941
```

```
62%|████████   | 31/50 [00:28<00:16, 1.14it/s]
```

```
Epoch 31/50 - Loss: 1.498 - Acc: 0.963
```

```
Val_loss: 1.515 - Val_acc: 0.945
```

```
82%|██████████ | 41/50 [00:37<00:08, 1.05it/s]
```

```
Epoch 41/50 - Loss: 1.493 - Acc: 0.969
```

```
Val_loss: 1.512 - Val_acc: 0.947
```

```
func:'train' took: 44.6958 sec
```

```
2%||          | 1/50 [00:00<00:40, 1.21it/s]
```

```
Epoch 1/50 - Loss: 1.500 - Acc: 0.960
```

```
Val_loss: 1.493 - Val_acc: 0.968
```

```
22%|██        | 11/50 [00:09<00:34, 1.14it/s]
```

```
Epoch 11/50 - Loss: 1.493 - Acc: 0.967
```

```
Val_loss: 1.496 - Val_acc: 0.964
```

```
42%|██████     | 21/50 [00:18<00:27, 1.05it/s]
```

```
Epoch 21/50 - Loss: 1.489 - Acc: 0.971
```

```
Val_loss: 1.498 - Val_acc: 0.962
```

```
62%|████████   | 31/50 [00:27<00:15, 1.23it/s]
```

```
Epoch 31/50 - Loss: 1.487 - Acc: 0.973
```

```
Val_loss: 1.498 - Val_acc: 0.962
```

```

82%|██████████| 41/50 [00:35<00:07, 1.25it/s]
Epoch 41/50 - Loss: 1.486 - Acc: 0.975
              Val_loss: 1.500 - Val_acc: 0.960

func:'train' took: 44.1680 sec

 2%||          | 1/50 [00:00<00:39, 1.24it/s]
Epoch 1/50 - Loss: 1.493 - Acc: 0.967
              Val_loss: 1.484 - Val_acc: 0.978

22%|██████    | 11/50 [00:08<00:31, 1.23it/s]
Epoch 11/50 - Loss: 1.488 - Acc: 0.971
              Val_loss: 1.485 - Val_acc: 0.976

42%|████████  | 21/50 [00:17<00:23, 1.23it/s]
Epoch 21/50 - Loss: 1.486 - Acc: 0.975
              Val_loss: 1.487 - Val_acc: 0.975

62%|██████████| 31/50 [00:26<00:16, 1.12it/s]
Epoch 31/50 - Loss: 1.485 - Acc: 0.975
              Val_loss: 1.489 - Val_acc: 0.971

82%|██████████| 41/50 [00:35<00:09, 1.03s/it]
Epoch 41/50 - Loss: 1.483 - Acc: 0.976
              Val_loss: 1.491 - Val_acc: 0.970

func:'train' took: 43.6335 sec

```

```
In [39]: ann_3_adam.evaluate(test_X_norm, test_y)
```

```

Accuracy: 0.9559999704360962
(tensor(1.5065), tensor(0.9560))

```

```
Out[39]:
```

(b) From the tabulated results, l2 regularization improved the accuracy a little (maybe not significant), but increased the run time by almost 2X.

Category	Training Accuracy	Validation Accuracy	Test Accuracy	Run Time(seconds)
1c: Hidden Layer=3	0.583	0.579	0.581	28.4
1d: Hidden Layer=50	0.986	0.983	0.961	42.1
2a: 1d + 15% dropout	0.976	0.970	0.956	43.6
2b: 1d + l2	0.988	0.988	0.962	78.9

From what I read, L2 regularization adds a penalty term to the loss function of the neural network to penalize the weights if they are too large. This method helps prevent overfitting by reducing the complexity of the model and promoting weight values that are more generalizable. But the step of adding regularization increases the training and computation complexity. The loss function with regularization now needs to be computed and backpropagated in each epoch / batch.



In our case, the benefit of adding L2 regularization doesn't seem very big. Likely because the ANN is not overfitting.

```
In [41]: ann_2_adam_l2 = Trainer(mlp2, "adam", 2e-3, 50, 128, input_transform=lambda x: x,)
```

```
# Split the dataset into 3-fold training and validation sets
training_result_all = []
kf = KFold(n_splits=3, shuffle=True, random_state=1)
for train_index, val_index in kf.split(train_X_norm, train_y):
    X_train, X_val = train_X_norm[train_index], train_X_norm[val_index]
    y_train, y_val = train_y[train_index], train_y[val_index]

    training_result = ann_2_adam_l2.train(X_train, y_train, X_val, y_val, \
        early_stop=False, l2=True, silent=False)
    training_result_all.append(training_result)
```

```
0%|          | 0/50 [00:00<?, ?it/s]<ipython-input-40-58e0eb1bad63>:137: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).
```

```
inputs = torch.tensor(inputs).float().clone().detach()
<ipython-input-40-58e0eb1bad63>:138: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).
```

```
outputs = torch.tensor(outputs).long().clone().detach()
```

```
2%||          | 1/50 [00:01<00:55, 1.12s/it]
```

```
Epoch 1/50 - Loss: 1.479 - Acc: 0.985
```

```
Val_loss: 1.472 - Val_acc: 0.987
```

```
22%|██        | 11/50 [00:10<00:36, 1.07it/s]
```

```
Epoch 11/50 - Loss: 1.474 - Acc: 0.987
```

```
Val_loss: 1.474 - Val_acc: 0.985
```

```
42%|████       | 21/50 [00:23<00:39, 1.35s/it]
```

```
Epoch 21/50 - Loss: 1.473 - Acc: 0.988
```

```
Val_loss: 1.476 - Val_acc: 0.984
```

```
62%|██████      | 31/50 [00:36<00:25, 1.33s/it]
```

```
Epoch 31/50 - Loss: 1.473 - Acc: 0.988
```

```
Val_loss: 1.479 - Val_acc: 0.983
```

```
82%|████████     | 41/50 [00:49<00:11, 1.31s/it]
```

```
Epoch 41/50 - Loss: 1.472 - Acc: 0.988
```

```
Val_loss: 1.481 - Val_acc: 0.982
```

```
func:'train' took: 62.0513 sec
```

```
2%||          | 1/50 [00:01<01:01, 1.26s/it]
```

```
Epoch 1/50 - Loss: 1.480 - Acc: 0.984
```

```
Val_loss: 1.473 - Val_acc: 0.988
```

```
22%|██        | 11/50 [00:14<00:51, 1.32s/it]
```

```
Epoch 11/50 - Loss: 1.473 - Acc: 0.988
```

```
Val_loss: 1.473 - Val_acc: 0.988
```

```
42%|████       | 21/50 [00:28<00:38, 1.33s/it]
```

```

Epoch 21/50 - Loss: 1.472 - Acc: 0.988
              Val_loss: 1.475 - Val_acc: 0.986
62%|██████    | 31/50 [00:45<00:38,  2.04s/it]
Epoch 31/50 - Loss: 1.472 - Acc: 0.989
              Val_loss: 1.477 - Val_acc: 0.985
82%|██████    | 41/50 [01:00<00:13,  1.56s/it]
Epoch 41/50 - Loss: 1.472 - Acc: 0.989
              Val_loss: 1.479 - Val_acc: 0.984

func:'train' took: 74.8374 sec
2%||          | 1/50 [00:01<01:10,  1.45s/it]
Epoch 1/50 - Loss: 1.479 - Acc: 0.984
              Val_loss: 1.470 - Val_acc: 0.990
22%|██        | 11/50 [00:16<00:55,  1.42s/it]
Epoch 11/50 - Loss: 1.473 - Acc: 0.988
              Val_loss: 1.470 - Val_acc: 0.990
42%|████      | 21/50 [00:31<00:44,  1.55s/it]
Epoch 21/50 - Loss: 1.473 - Acc: 0.988
              Val_loss: 1.471 - Val_acc: 0.989
62%|██████    | 31/50 [00:46<00:30,  1.59s/it]
Epoch 31/50 - Loss: 1.473 - Acc: 0.988
              Val_loss: 1.473 - Val_acc: 0.989
82%|██████    | 41/50 [01:03<00:15,  1.67s/it]
Epoch 41/50 - Loss: 1.472 - Acc: 0.988
              Val_loss: 1.475 - Val_acc: 0.988

```

```
func:'train' took: 78.8620 sec
```

```
In [43]: ann_2_adam_l2.evaluate(test_X_norm, test_y)
```

```
Accuracy: 0.9623000025749207
(tensor(1.5031), tensor(0.9623))
```

```
Out[43]:
```

(c) I used pca to transform the Train\_X\_norm dataset with 99% variance kept. Then I fit the pca model to the test\_X\_norm dataset. The features decreased from 1024 to 331. I modified the ANN model and changed the 1st layer from 1024 to 331 then ran the training.

The results are summarized below:

Category	Training Accuracy	Validation Accuracy	Test Accuracy	Run Time(seconds)
1c: Hidden Layer=3	0.583	0.579	0.581	28.4
1d: Hidden Layer=50	0.986	0.983	0.961	42.1
2a: 1d + 15% dropout	0.976	0.970	0.956	43.6

Category	Training Accuracy	Validation Accuracy	Test Accuracy	Run Time(seconds)
2b: 1d + l2	0.988	0.988	0.962	78.9
2c: 1d + pca	0.981	0.980	0.955	24.2

Overall the pca method works almost as well as the original ANN from 1d. The time is much shorter now because we have less features to fit. The accuracies are comparable. I think it's a good data transformation technique.

```
In [59]: pca = PCA(0.99)
train_X_pca = pca.fit_transform(train_X_norm.reshape(60000, 1024))
test_X_pca = pca.transform(test_X_norm.reshape(10000, 1024))
print("train_X shape after PCA transformation: ", train_X_pca.shape)
print("test_X shape after PCA transformation: ", test_X_pca.shape)
```

```
train_X shape after PCA transformation: (60000, 331)
test_X shape after PCA transformation: (10000, 331)
```

```
In [61]: class MLP2_pca(nn.Module):
    def __init__(self):
        super(MLP2_pca, self).__init__()
        self.layers = nn.Sequential(
            nn.Flatten(),
            nn.Linear(331, 50),
            nn.Sigmoid(),
            nn.Linear(50, 10),
            nn.Sigmoid()
        )

    def forward(self, X):
        return self.layers(X)
```

```
In [62]: MLP2_pca = MLP2_pca()
ann_2_adam_pca = Trainer(MLP2_pca, "adam", 2e-3, 50, 128, input_transform=lambda x: x,)

# Split the dataset into 3-fold training and validation sets
training_result_all = []
kf = KFold(n_splits=3, shuffle=True, random_state=1)
for train_index, val_index in kf.split(train_X_pca, train_y):
    X_train, X_val = train_X_pca[train_index], train_X_pca[val_index]
    y_train, y_val = train_y[train_index], train_y[val_index]

    training_result = ann_2_adam_pca.train(X_train, y_train, X_val, y_val, \
        early_stop=False, l2=False, silent=False)
    training_result_all.append(training_result)
```

```
0%|          | 0/50 [00:00<?, ?it/s]<ipython-input-40-58e0eblbad63>:137: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).
  inputs = torch.tensor(inputs).float().clone().detach()
<ipython-input-40-58e0eblbad63>:138: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).
  outputs = torch.tensor(outputs).long().clone().detach()
2%||         | 1/50 [00:00<00:21, 2.25it/s]
Epoch 1/50 - Loss: 1.933 - Acc: 0.774
              Val_loss: 1.688 - Val_acc: 0.887

22%|█        | 11/50 [00:05<00:17, 2.17it/s]
Epoch 11/50 - Loss: 1.518 - Acc: 0.947
              Val_loss: 1.531 - Val_acc: 0.933

42%|██       | 21/50 [00:09<00:12, 2.35it/s]
Epoch 21/50 - Loss: 1.498 - Acc: 0.962
              Val_loss: 1.521 - Val_acc: 0.941

62%|████    | 31/50 [00:14<00:08, 2.19it/s]
Epoch 31/50 - Loss: 1.489 - Acc: 0.970
              Val_loss: 1.518 - Val_acc: 0.944

82%|██████   | 41/50 [00:19<00:03, 2.26it/s]
Epoch 41/50 - Loss: 1.483 - Acc: 0.975
              Val_loss: 1.518 - Val_acc: 0.945

func:'train' took: 22.9156 sec

2%||         | 1/50 [00:00<00:21, 2.31it/s]
Epoch 1/50 - Loss: 1.497 - Acc: 0.963
              Val_loss: 1.484 - Val_acc: 0.976

22%|█        | 11/50 [00:05<00:21, 1.85it/s]
Epoch 11/50 - Loss: 1.483 - Acc: 0.975
              Val_loss: 1.489 - Val_acc: 0.973

42%|██       | 21/50 [00:09<00:12, 2.33it/s]
Epoch 21/50 - Loss: 1.479 - Acc: 0.978
              Val_loss: 1.495 - Val_acc: 0.969

62%|████    | 31/50 [00:13<00:07, 2.38it/s]
Epoch 31/50 - Loss: 1.477 - Acc: 0.980
              Val_loss: 1.499 - Val_acc: 0.966

82%|██████   | 41/50 [00:18<00:04, 1.91it/s]
Epoch 41/50 - Loss: 1.476 - Acc: 0.981
              Val_loss: 1.502 - Val_acc: 0.964

func:'train' took: 22.9999 sec

2%||         | 1/50 [00:00<00:21, 2.25it/s]
Epoch 1/50 - Loss: 1.490 - Acc: 0.970
              Val_loss: 1.474 - Val_acc: 0.986

22%|█        | 11/50 [00:04<00:16, 2.37it/s]
Epoch 11/50 - Loss: 1.479 - Acc: 0.978
              Val_loss: 1.475 - Val_acc: 0.985
```

```

42%|███████| 21/50 [00:11<00:13, 2.13it/s]
Epoch 21/50 - Loss: 1.477 - Acc: 0.979
              Val_loss: 1.479 - Val_acc: 0.983

62%|███████| 31/50 [00:15<00:08, 2.35it/s]
Epoch 31/50 - Loss: 1.476 - Acc: 0.981
              Val_loss: 1.483 - Val_acc: 0.982

82%|███████| 41/50 [00:20<00:04, 1.82it/s]
Epoch 41/50 - Loss: 1.476 - Acc: 0.981
              Val_loss: 1.486 - Val_acc: 0.980

func:'train' took: 24.2085 sec

```

```
In [63]: ann_2_adam_pca.evaluate(test_X_pca, test_y)
```

```

Accuracy: 0.9545000195503235
(tensor(1.5128), tensor(0.9545))

```

```
Out[63]:
```

(d) I chose L2 regularization + PCA transformation. The results are summarized below:

Category	Training Accuracy	Validation Accuracy	Test Accuracy	Run Time(seconds)
1c: Hidden Layer=3	0.583	0.579	0.581	28.4
1d: Hidden Layer=50	0.986	0.983	0.961	42.1
2a: 1d + 15% dropout	0.976	0.970	0.956	43.6
2b: 1d + l2	0.988	0.988	0.962	78.9
2c: 1d + pca	0.981	0.980	0.955	24.2
2d: 1d + l2 + pca	0.986	0.9802	0.958	28.6

Overall the combination of L2 regularization and PCA data transformation gives very good results. The training is faster than both 1d and 2b cases. The accuracies are also very decent. In general this approach combines the benefits of both L2 regularization and the PCA transformation.

```

In [66]: ann_2_adam_l2_pca = Trainer(MLP2_pca, "adam", 2e-3, 50, 128, input_transform=lambda x: x,)

# Split the dataset into 3-fold training and validation sets
training_result_all = []
kf = KFold(n_splits=3, shuffle=True, random_state=1)
for train_index, val_index in kf.split(train_X_pca, train_y):
    X_train, X_val = train_X_pca[train_index], train_X_pca[val_index]
    y_train, y_val = train_y[train_index], train_y[val_index]

    training_result = ann_2_adam_l2_pca.train(X_train, y_train, X_val, y_val, \

```

```
early_stop=False, l2=True, silent=False)
training_result_all.append(training_result)
```

```
0%|          | 0/50 [00:00<?, ?it/s]<ipython-input-40-58e0eb1bad63>:137: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).
```

```
inputs = torch.tensor(inputs).float().clone().detach()
<ipython-input-40-58e0eb1bad63>:138: UserWarning: To copy construct from a tensor, it is recommended to use sourceTensor.clone().detach() or sourceTensor.clone().detach().requires_grad_(True), rather than torch.tensor(sourceTensor).
```

```
outputs = torch.tensor(outputs).long().clone().detach()
```

```
2%||          | 1/50 [00:00<00:30, 1.60it/s]
```

```
Epoch 1/50 - Loss: 1.482 - Acc: 0.981
```

```
Val_loss: 1.474 - Val_acc: 0.984
```

```
22%|██        | 11/50 [00:07<00:22, 1.77it/s]
```

```
Epoch 11/50 - Loss: 1.475 - Acc: 0.985
```

```
Val_loss: 1.475 - Val_acc: 0.984
```

```
42%|████      | 21/50 [00:12<00:15, 1.85it/s]
```

```
Epoch 21/50 - Loss: 1.474 - Acc: 0.986
```

```
Val_loss: 1.479 - Val_acc: 0.982
```

```
62%|██████    | 31/50 [00:18<00:10, 1.78it/s]
```

```
Epoch 31/50 - Loss: 1.474 - Acc: 0.986
```

```
Val_loss: 1.484 - Val_acc: 0.979
```

```
82%|████████  | 41/50 [00:24<00:05, 1.75it/s]
```

```
Epoch 41/50 - Loss: 1.474 - Acc: 0.986
```

```
Val_loss: 1.488 - Val_acc: 0.977
```

```
func:'train' took: 30.1707 sec
```

```
2%||          | 1/50 [00:00<00:27, 1.81it/s]
```

```
Epoch 1/50 - Loss: 1.483 - Acc: 0.981
```

```
Val_loss: 1.475 - Val_acc: 0.985
```

```
22%|██        | 11/50 [00:06<00:21, 1.78it/s]
```

```
Epoch 11/50 - Loss: 1.475 - Acc: 0.985
```

```
Val_loss: 1.475 - Val_acc: 0.985
```

```
42%|████      | 21/50 [00:12<00:17, 1.66it/s]
```

```
Epoch 21/50 - Loss: 1.474 - Acc: 0.986
```

```
Val_loss: 1.478 - Val_acc: 0.984
```

```
62%|██████    | 31/50 [00:18<00:10, 1.74it/s]
```

```
Epoch 31/50 - Loss: 1.474 - Acc: 0.986
```

```
Val_loss: 1.483 - Val_acc: 0.981
```

```
82%|████████  | 41/50 [00:24<00:05, 1.65it/s]
```

```
Epoch 41/50 - Loss: 1.474 - Acc: 0.987
```

```
Val_loss: 1.487 - Val_acc: 0.978
```

```
func:'train' took: 29.7061 sec
```

```
2%||          | 1/50 [00:00<00:25, 1.89it/s]
```

```
Epoch 1/50 - Loss: 1.483 - Acc: 0.980
```

```
Val_loss: 1.472 - Val_acc: 0.989
```

```
22%|██████| 11/50 [00:06<00:23, 1.66it/s]
Epoch 11/50 - Loss: 1.475 - Acc: 0.985
              Val_loss: 1.472 - Val_acc: 0.989

42%|███████| 21/50 [00:12<00:15, 1.87it/s]
Epoch 21/50 - Loss: 1.475 - Acc: 0.985
              Val_loss: 1.476 - Val_acc: 0.987

62%|████████| 31/50 [00:18<00:12, 1.49it/s]
Epoch 31/50 - Loss: 1.475 - Acc: 0.986
              Val_loss: 1.480 - Val_acc: 0.984

82%|█████████| 41/50 [00:23<00:04, 1.86it/s]
Epoch 41/50 - Loss: 1.474 - Acc: 0.986
              Val_loss: 1.484 - Val_acc: 0.982
```

```
func:'train' took: 28.6386 sec
```

```
In [68]: ann_2_adam_l2_pca.evaluate(test_X_pca, test_y)
```

```
Accuracy: 0.9577000141143799
Out[68]: (tensor(1.5090), tensor(0.9577))
```